

Prob and Stat with R

Roudranil Das

PBSR Aug-Nov 2022

Roll No: MDS202227

roudranil@cmi.ac.in

October 17, 2022

Contents

1	August 08, 2022	2
1.1	Random sampling and randomised controlled clinical trials	2
1.2	Sampling from social media data	2
2	August 10, 2022	2
2.1	Social media data cant be used	2
2.2	Hands on with R	2
3	August 17, 2022	3
3.1	Randomization	3
4	August 22, 2022	3
4.1	Random Variable	3
4.1.1	Continuous random variable	3
5	August 24, 2022	4
5.1	Continuous Random Variable (contd.)	4
5.2	Moments	4
6	September 1, 2022	4
6.1	Central moments	4
6.2	Moment generating function	5
6.3	Conditional Probability	5
7	September 5, 2022	5
7.1	SRSWR and SRSWOR	5
8	8 September, 2022	5
8.1	Mixture models	5
9	12 September, 2022	7
9.1	Datasets analysed	7
9.2	Sampling distributions	7
9.3	Distribution under transformation	8
10	14 September, 2022	8
10.1	Central Limit Theorem	8
10.2	Probability distributions under transformations	9

11 19 September, 2022	9
11.1 Invariance property of MLE	9
11.2 Transformation of random variables	9
12 21 September, 2022	10
12.1 Some exercises	10
12.2 Joint Probability	11
13 22 September, 2022	11
13.1 Joint prob cotd	11
14 17 October, 2022	12
14.1 Linear Regression cotd	12
14.2 Statistical Inference	12

1 August 08, 2022

1.1 Random sampling and randomised controlled clinical trials

If clinical trials are not randomised, then what problems are there?

Problems are mainly ethical or there might be a judgemental bias in choosing which participant is in which group.

Ideally we would want all data/population groups to be proportionately or equally represented in all groups. This is impossible to do by ourselves. In this, Fisher's suggestion was to assign them by tossing an unbiased coin. By the time we are done, the proportions are taken care of. We can also do the trials across multiple phases.

1.2 Sampling from social media data

Social media is **voluntary response survey methodology** (which always gives you an extremely biased view of the situation.) For good quality of data we should survey it properly, and although randomised survey methodology is generally expensive, the data it provides is usually much more representative of the population.

2 August 10, 2022

2.1 Social media data cant be used

In last class, we had seen how social media data cannot be used in order to predict political opinion (check **"The Great Hack"**). This is because this data is generally **oversampled**, that is some subset of the population is overrepresented in this population.

2.2 Hands on with R

Check the R markdown file. Some commands and stuff about R:

- While indexing, putting a negative integer removes that entry.
- `rbind.data.frame` binds rows of a data frame to an existing data frame.
-

3 August 17, 2022

3.1 Randomization

To understand the concept of randomization, note the issue in clinical trials again. We need randomised assignment of participants in the two groups. Check the R markdown file. Say we have created two groups of control and experiment, and both samples have very similar characteristics across the board. Now if the experiment group shows vastly different results compared to the control group, we will know that whatever it is that we were testing has a really good effect as there was no inherent bias in how the groups were formed.

4 August 22, 2022

4.1 Random Variable

What is a random variable?

Usual definitions: function from the sample space to \mathbb{R} . Example: an accident is a random event. Following that we can have the following random variables:

- (i) dead or alive (binary, discrete)
- (ii) bill amount (continuous)
- (iii) number of days in rehab (discrete, non binary)

$$X : \Omega \longrightarrow \mathbb{R} \quad (1)$$

We are associated in probability associated with random variable. Convention: $A, B, C, D \rightarrow \text{Random events}$, $X, Y, Z, W \rightarrow \text{Random variables}$. If they are in small letters, they are either observed events or realised random variables.

Say we want to try and introduce the idea of convergence in this probability. So we would need $0 < p(x) < 1$. Moreover since the sum of all probabilities should be 1, the infinite series with sum of probabilities should be convergent. Say we are dealing with discrete random variables. Then

$$P[X \leq x] = \sum_{n=0}^x p(n) \rightarrow c < \infty \text{ as } x \rightarrow \infty$$

Here this following sequence will be very handy:

$$S(n) = \frac{\lambda^n}{n!}$$

Note that

$$\sum_{n=0}^{\infty} S(n) = e^{\lambda}$$

We define, $p(x) = \frac{S(x)}{e^{\lambda}} = e^{-\lambda} \frac{\lambda^x}{x!}$. Thus we have a probability distribution and the corresponding pmf (of **Poisson** distribution).

$$P[X \leq x] = \sum_{n=0}^x e^{-\lambda} \frac{\lambda^n}{n!} \quad (2)$$

Cumulative distribution function \uparrow

4.1.1 Continuous random variable

For probability distribution for continuous r.v.,

$$P[X \leq x] = \int_{-\infty}^x f(x) dx$$

5 August 24, 2022

5.1 Continuous Random Variable (contd.)

From Fisher's Definition, if we have data in bins and relative frequency, then area of bins is relative frequency which is an approximation of the probability. Suppose we have a histogram, y-axis being **density** and x-axis being the bins. then $P[c < X < d] = B_1 + B_2 + B_3 + B_4 = \sum_{i=1}^4 B_i = \sum_{i=1}^4 d_i \times \Delta_i$. Thus we will sort of need to approximate the density d_i by some function $f(x_i)$. And we will need $f(x_i) > 0$.

Thus,

$$\begin{aligned} P[a < X < b] &= \sum_{i=1}^k f(x_i)(x_i - x_{i-1}) \\ &= \int_a^b f(x)dx \text{ (as } \Delta \rightarrow 0) \end{aligned}$$

This is the probability density function, with the usual properties. Thus by mathematical construction, $P[X = a] = 0$.

$F_X(a) = P[X < a]$ (cdf).

5.2 Moments

Say from a distribution, we have x_1, x_2, \dots, x_n (SRSWR). Then define, $\bar{x} = \frac{\sum x_i}{n}$ (**sample average**). We can sort of rewrite this as $\bar{x} = \sum x_i w_i$ where $w_i = \frac{1}{n} \forall i$. Then

$$\bar{x} = \sum x_i p(x_i) = \mathbb{E}[X]$$


For continuous,

$$\mathbb{E}[X] = \int x f(x) dx$$

This is **First order raw moments**. **N^{th} order raw moments** will be $\mathbb{E}[X^n]$.

6 September 1, 2022

6.1 Central moments



$$\mathbb{V}(X) = \int (x - \mu)^2 f(x) dx = \mathbb{E}[x - \mu]^2 \quad (3)$$

This is sort of a probability version. For sample, if the random sample drawn is $\{x_1, x_2, \dots, x_n\}$ then the sample mean is $\bar{x} = \frac{1}{n} \cdot \sum x_i$. Then the sample variance is $s^2 = \frac{1}{n} \cdot \sum (x_i - \bar{x})^2$. Or if we take the unbiased estimator for $s^2 = \frac{1}{n-1} \cdot \sum (x_i - \bar{x})^2$.

6.2 Moment generating function

$\mathbb{E}[e^{tX}] = \int e^{tX} f(x) dx$ is the moment generating function. Specifically, $\frac{\partial \mathbb{E}[e^{tX}]}{\partial t} = \mathbb{E}[X]$ at $t = 0$.

Why? Because by expanding $e^{tX} = 1 + tX + \frac{(tX)^2}{2!} + \dots$ and taking expectation we get $\mathbb{E}[e^{tX}] = \mathbb{E}[1 + tX + \frac{(tX)^2}{2!} + \dots]$. Hence after differentiating wrt t and setting $t = 0$, we get the expectation. Higher order partial derivatives give the other moments.

6.3 Conditional Probability

Say we have two events A and S . Then

- $\mathbb{P}(A) \rightarrow$ marginal probability of A
- $\mathbb{P}(A | S) \rightarrow$ conditional probability of A given S

If $\mathbb{P}(A) = \mathbb{P}(A | S)$ then the two events are independent.

For random variables, $\mathbb{P}(X) \rightarrow$ prob dist of X . $\mathbb{P}(X | Y = y) \rightarrow$ cond prob dist of X given $Y = y$. $\mathbb{P}(X, Y) \rightarrow$ Joint prob dist. If $\mathbb{P}(X, Y) = \mathbb{P}(X) \cdot \mathbb{P}(Y)$, then the two random variables are independent.

7 September 5, 2022

7.1 SRSWR and SRSWOR

Say we have an urn, with multiple balls of different colours. Then say we draw a sample of size 3, (SRSWR), and record the colours. Also we do this SRSWOR (another sample of size = 3).

- In the first draw if we see blue, what is the probability we will see a blue ball in the next draw?

$$\mathbb{P}(B_2 | B_1) = \begin{cases} \frac{1}{5}, \text{SRSWR} \\ 0, \text{SRSWOR} \end{cases}$$

In SRSWR, **draws are independent**. In SRSWOR **draws are not independent**.


There are two kinds of data set we will need to work with in order to understand the two simple random sampling schemes.

1. Cross-sectional Data set: Essentially SRSWR. If we have an experiment going on for a long time, then at one time point, the data across all the attributes is cross-sectional data set.
2. Time-series Data set (or transversal data set): Series of data indexed by time points. Weather data, profit and loss data across a long time. And because information from today affects information of the future, the observations are not independent. Hence, the data is SRSWOR.
3. Longitudinal Data set (panel data): Track the same sample at different points of time. For any one sample, the data will be like a time series. But at any time point, the samples will behave like a cross-sectional data set.


8 8 September, 2022

8.1 Mixture models

We take multiple probability distributions, then their convex combination is another probability distribution. This can be used to approximate complex distributions with different properties. For Example:



This was approximated by:



This is $p \cdot N(52, 5.5^2) + (1 - p) \cdot N(80, 6.5^2)$, where $0 < p < 1$.

Say we want to compute the MLE of parameters of a bimodal mixed distribution, like this one. Then we define the negative log likelihood and use the `optim` function to find the optimum parameters.

```
1 NegLogLikeMix <- function(theta, data){
2   mu1 = theta[1]
3   sigma1 = exp(theta[2]) # because optim runs it on the real line, but sigma needs to
   be positive
```

```

4      mu2 = theta[3]
5      sigma2 = exp(theta[4])
6
7
8      p = exp(theta[5])/(1+exp(theta[5])) # logit transformation
9
10     n = length(data)
11
12     l = 0
13     for(i in 1:n){
14         l = l + log(p*dnorm(data[i], mean = mu1, sd = sigma1) + (1-p)*dnorm(data[i],
15                                     mean = mu2, sd = sigma2))
16     }
17     return(-l)
18 }
19 theta_initial = c(52,8,80,8,0.5)
20 optim(theta_initial,
21       NegLogLikMix,
22       data=waiting,
23       control = list(maxit = 1500))

```

What we are doing here is simple likelihood optimisation.

Say now we have the optimal parameters. So we say the model has been fit. Now what?

9 12 September, 2022

9.1 Datasets analysed

1. Football Dataset :-
 - Poisson model
 - (a) method of moments
 - (b) Maximum likelihood estimation to find the unknown parameter of the model λ
 - Using Poisson Model we did some inference.
2. Old faithful geyser dataset :-
 - $(\begin{smallmatrix} x_1 \\ y_1 \end{smallmatrix}), \dots, (\begin{smallmatrix} x_n \\ y_n \end{smallmatrix})$ data points - We had a mixture model.
 - (a) used the likelihood method to estimate the 5 parameters (check the rmd)

Estimate $\hat{\theta}$ is not going to be exactly θ_0 (true parameter). We don't actually know this θ_0 . We have a dataset $D = \{X_1, X_2, \dots, X_n\}$ and it contains some information about θ_0 .

9.2 Sampling distributions

Suppose we are doing a survey in order to find out the avg income of an area. So we take random samples and calculate their sample means (given that we have ∞ resources).

1. $[x_1, x_2, \dots, x_n] = \tilde{x}_1, \hat{x}_1$
2. $[x_1, x_2, \dots, x_n] = \tilde{x}_2, \hat{x}_2$
- .
- .
- .
1000. $[x_1, x_2, \dots, x_n] = \tilde{x}_{1000}, \hat{x}_{1000}$

Say we make an histogram with \hat{x}_r . Then we will get an idea of their probability distribution. This is the **sampling distribution of the sample means**.

Sampling distribution of an estimator (not necessarily sample mean) tells us that what is the expected error that estimator can make, and what is the unlikely error for that estimator.

Say for a bell shaped distribution, mean = median. Then

Q: Shall we use sample mean or sample median?

We build the sampling distribution of both, and take the one with less standard error. Or its better if we make an educated choice.

When we increase the sample size the sampling distribution becomes sharper towards the true value (at the limiting case it will become a degenerate distribution). And if we take a small distribution, then it will become flatter.

9.3 Distribution under transformation

(read CSBGR chap2 pgs 47 - 68)

Exercise 9.3.1. $X \sim N(\mu, \sigma^2)$. Calculate the moment generating function of X .

$$M_X(t) = \mathbb{E}[e^{tX}] = \int_{-\infty}^{\infty} f(x)dx = e^{\mu t + \frac{1}{2}\sigma^2 t^2}$$

Now say $X_1, X_2, \dots, X_n \stackrel{iid}{\sim} N(\mu, \sigma^2)$, $S = \sum X_i$, $\bar{X} = \frac{S}{n}$. Then what is the distribution of S ? To find that we find $M_S(t)$.

$$M_S(t) = \mathbb{E}[e^{tS}] = \mathbb{E}\left[e^{t \sum X_i}\right]$$

$$= \mathbb{E}\left[\prod e^{tX_i}\right] = \prod \mathbb{E}\left[e^{tX_i}\right]$$

We could write the last line because of the iid. From there we can say that $S \sim N(n\mu, n\sigma^2)$.

Exercise 9.3.2. Show $\bar{X} \stackrel{iid}{\sim} N(\mu, \frac{\sigma^2}{n})$

Proof. Method 1: Using the independence.

Given that $X_1, X_2, \dots, X_n \stackrel{iid}{\sim} N(\mu, \sigma^2)$,

$$\mathbb{E}[\bar{X}_n] = \frac{1}{n} \mathbb{E}[\sum X_i] = \mu$$

$$\mathbb{V}(\bar{X}_n) = \frac{1}{n^2} \sum \mathbb{V}(X_i) = \frac{\sigma^2 n}{n^2} = \frac{\sigma^2}{n}.$$

$$\therefore \bar{X} \sim N(\mu, \frac{\sigma^2}{n})$$

Better method would be with mgf. □

10 14 September, 2022

10.1 Central Limit Theorem

Lindeberg-Levy's Central Limit Theorem (CLT)

Suppose $\{X_1, X_2, \dots, X_n\}$ are iid random variables such that:

- $\mathbb{E}[X_i] = \mu$
- $\mathbb{V}(X_i) = \sigma^2 < \infty$ (automatically this means that the mean exists)

Then as $n \rightarrow \infty$,

$$\frac{\sqrt{n}(\bar{X} - \mu)}{\sigma} \xrightarrow{\mathcal{L}} N(0, 1)$$

What this means is that, if we have large enough sample size, $\bar{X} \approx N(\mu, \frac{\sigma^2}{n})$. Most importantly, this theorem does not assume any sort of underlying distribution. We only need their mean and variance to exist.

Sampling distribution of sample proportion will approximately follow normal if we can ensure large enough sample sizes.

$$\hat{p} = \bar{X} \overset{approx}{\sim} N\left(p, \frac{p(1-p)}{n}\right)$$

Previously we checked if the underlying distribution is Gamma, then the sampling distribution of the sample mean is also Gamma. But due to CLT, as sample size grows, the exact sampling distribution which is Gamma will start behaving like Gaussian.

Typically we call this $P_n \xrightarrow{\mathcal{L}} P$ or $P_n \xrightarrow{\mathcal{D}} P$: convergence in law or convergence in distribution. CLT only talks about sampling distribution of sample mean.

10.2 Probability distributions under transformations

If X is a random variable with cdf $F_X(x)$, $Y = g(X)$ and g is deterministic. Then Y is also a random variable.

Example. X is rv. $Y = \log(X)$ is a rv.

$$P(Y \in A) = P(g(X) \in A)$$

Say, $X \in \mathcal{X}$: sample space (range of values of X). Then,

$$g(x) : \mathcal{X} \longrightarrow \mathcal{Y}$$

We assume that inverse mapping of g also exists and we denote it as g^{-1} . For $A \subseteq \mathcal{Y}$

$$g^{-1}(A) = \{x \in \mathcal{X} : g(x) \in A\}$$

The 3 axioms of Kolmogorov holds for $P(Y \in A)$ under the transformation $Y = g(X)$.

11 19 September, 2022

11.1 Invariance property of MLE

Invariance property

if mle of θ is $\hat{\theta}$ then mle of $\tau(\theta)$ is $\tau(\hat{\theta})$

Example: $X_1, X_2, \dots, X_n \overset{iid}{\sim} \text{Bernoulli}(p)$ $X_i = 1$ if i^{th} person develops some disease. Then $\frac{p}{1-p}$ is the odds. Sample proportion $\frac{\sum X_i}{n}$ is the MLE of p . So we can use invariance property here.

Now $\hat{\theta}$ has its own sampling distribution $\hat{\theta} \sim f_{\hat{\theta}}$. Does $\tau(\hat{\theta})$ has its own sampling distribution? Yes it does by last day's transformation result.

11.2 Transformation of random variables

Suppose $X \sim \text{Bin}(n, p)$. Define $Y = n - X$. Then we can say $Y \sim \text{Bin}(n, q)$ where $q = 1 - p$. To do this, $P[Y = a] = P[X = n - a] = \binom{n}{n-a} p^{n-a} q^a$. But, $P[Y = a] = \binom{n}{a} q^a p^{n-a}$. Thus $Y \sim \text{Bin}(n, q)$.

So, say we have $Y = g(X)$. Then we assume g to be monotone, ie, one-one and onto. When this condition holds its actually much easier to find out the distribution of the transformed random variable.

Exercise 11.2.1. $X \sim \text{Unif}(0, 2\pi)$. Then $f_X(x) = \frac{1}{2\pi}$ when $x \in (0, 2\pi)$, 0 otherwise. Consider the transformation $Y = \sin^2(x)$.

Theorem

Suppose X has a CDF $F_X(x)$ and pdf $f_X(x)$.

$$\mathcal{X} = \{x : f_X(x) > 0\}$$

$$\mathcal{Y} = \{y : y = g(x) \text{ for some } x \in \mathcal{X}\}$$

1. if g is increasing function, then $F_Y(y) = F_X(g^{-1}(y))$
2. if g is decreasing function, then $F_Y(y) = 1 - F_X(g^{-1}(y))$

Theorem

Let X have pdf $f_X(x)$ and let $Y = g(X)$ where g is monotone. Let \mathcal{X} and \mathcal{Y} be defined as above. Suppose $f_X(x)$ is continuous on \mathcal{X} . $g^{-1}(y)$ has a continuous derivative on \mathcal{Y} . Then the pdf of Y is given by:

$$f_Y(y) = f_X(g^{-1}(y)) \left| \frac{d}{dy} g^{-1}(y) \right|, y \in \mathcal{Y} \quad (4)$$

Even though we work very much with the normal distribution, it's cdf does not have an analytical closed form solution. It is always solved using numerical methods.

12 21 September, 2022

12.1 Some exercises

Exercise 12.1.1 (Inverted Gamma, C.B. pg 51). The pdf of the Gamma distribution is as follows:

$$f(x) = \begin{cases} \frac{1}{\Gamma(n)\beta^n} x^{n-1} e^{-\frac{x}{\beta}}, 0 < x < \infty \\ 0, \text{otherwise} \end{cases}$$

β : positive constant, $n \in \mathbb{N}$

$$Y = g(X) = \frac{1}{X} \Rightarrow g^{-1}(y) = \frac{1}{y}$$

Proof. Then we have the Jacobian, $J = \frac{d}{dy} g^{-1}(y) = -\frac{1}{y^2}$. Therefore

$$f_Y(y) = \frac{1}{\Gamma(n)\beta^n} \left(\frac{1}{y} \right)^{n+1} e^{-\frac{1}{\beta \cdot y}}$$

□

Exercise 12.1.2 (C.B. pg 52). $X \sim$ continuous rv. Consider the transformation $Y = X^2$. What will be the CDF of Y ?

Proof.

$$\begin{aligned} F_Y(y) &= P[Y \leq y] \\ &= P[X^2 \leq y] \\ &= P[X \leq \sqrt{y}] - P[X \leq -\sqrt{y}] \\ &= F_X(\sqrt{y}) - F_X(-\sqrt{y}) \end{aligned}$$

For pdf, $f_Y(y) = \frac{d}{dy} F_Y(y)$

□

Exercise 12.1.3. $X \sim N(0, 1)$

$$f(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}}, -\infty < x < \infty$$

And we have been given that $Y = X^2$. Calculate the pdf.

Proof. We know that the transformation is monotone in $(-\infty, 0)$ and $(0, \infty)$. Using the previous result,

$$f_Y(y) = \frac{1}{\sqrt{2\pi}} e^{\frac{-(-\sqrt{y})^2}{2}} \left| \frac{-1}{2\sqrt{y}} \right| + \frac{1}{\sqrt{2\pi}} e^{\frac{-(-\sqrt{y})^2}{2}} \left| \frac{1}{2\sqrt{y}} \right| + f(0)|J|$$

$$\Rightarrow f_Y(y) = \frac{1}{\sqrt{2\pi}} \frac{1}{\sqrt{y}} e^{-y/2}, y > 0$$

Actually, $Y \sim \chi_{(1)}^2 = \text{Gamma}(\alpha=?, \beta=?)$ □

Actually if we have $Y_1, \dots, Y_n \stackrel{iid}{\sim} \chi_{(1)}^2 \equiv \text{Gamma}(\alpha=?, \beta=?)$, then we can show, $\sum Y_i \sim \chi_{(n)}^2$.

12.2 Joint Probability

Say we have two events, A, B . Then the joint probability of them is $P(A \cap B)$.

Example. A : children with **age** < 6 , B : children with **BMI** < 20

Then a statistician will obviously make some sort of plot of BMI vs age . Say it is a scatter plot with univariate histograms on the margin. Then we can find the area where both the conditions hold. The volume of the cuboid then gives us the relative frequency (much later, the probability).

$P[x_i \leq X \leq x_{i+1}, y_i \leq Y \leq y_{i+1}] \approx r = \frac{f}{n}$. To get the density, $\Delta x_i = x_{i+1} - x_i$. Similarly, Δy_i . Let $d = \frac{r}{\Delta x_i \Delta y_i}$. Then,

$$P[x_i \leq X \leq x_{i+1}, y_i \leq Y \leq y_{i+1}] = d \times \Delta x_i \Delta y_i$$

- Volume of the tower is relative frequency
- Height of the tower is the density
- Area of the box represents the set for which we are interested in finding the probability.

We can expand this concept to a box between any 4 points: $P[a \leq X \leq b, c \leq Y \leq d]$. This we can do by double summation (for discrete) or by double integration (for continuous).

13 22 September, 2022

13.1 Joint prob cotd

We want to assume $d \sim$ smooth surface, say, $d \sim f(x, y)$ such that

1. $f(x, y) \geq 0, x, y \in \mathbb{R}^2$
2. $\int \int_{\mathbb{R}^2} f(x, y) dx dy = 1$

This is the joint probability density function of the two variables.

To get the pdf of the marginal distribution of individual variables,

$$\int_{y \in \mathbb{R}} f(x, y) dy = f(x)$$

or,

$$\int_{x \in \mathbb{R}} f(x, y) dx = f(y)$$

14 17 October, 2022

14.1 Linear Regression ctd

$$\mathbb{E}[Y | X] = \mu_y + \rho \frac{\sigma_y}{\sigma_x} (x - \mu_x)$$

The above is the linear regression of Y on X . We can rewrite this as $E[Y | X] = \alpha + \beta \cdot x$ where $\alpha = \left(\mu_y - \rho \frac{\sigma_y}{\sigma_x} \mu_x\right)$ and $\beta = \rho \frac{\sigma_y}{\sigma_x}$. From the data we can easily get $\mu_x = \bar{x}, \mu_y = \bar{y}, \sigma_x = s_x, \sigma_y = s_y$. Then,

$$\rho : r = \frac{s_{xy}}{s_x s_y} \text{ where } s_{xy} = \frac{1}{n} \sum (x_i - \bar{x})(y_i - \bar{y})$$

We have $(x_1, y_2) \dots (x_n, y_n) \sim \text{BN}(\dots)$. But also,

$$X_1, \dots, X_n \sim N(\mu_x, \sigma_x)$$

$$Y_1, \dots, Y_n \sim N(\mu_y, \sigma_y)$$

$$\bar{X} \sim N\left(\mu_x, \frac{\sigma_x^2}{n}\right), \bar{Y} \sim N\left(\mu_y, \frac{\sigma_y^2}{n}\right)$$

$$\frac{ns_x^2}{\sigma_x^2} \sim \chi_{df}^2, \frac{ns_y^2}{\sigma_y^2} \sim \chi_{df}^2$$

Can we figure out the distributions of $\hat{\alpha}$ and $\hat{\beta}$?

In the above paragraph, we have actually used the maximum likelihood estimators, shouldnt we use the unbiased estimator?

Each of this estimators have their own probability distributions, so we shouldnt just stop at finding the estimated values. Next up we have to give a margin of error for our estimates. We can only get this if we internalise the idea that these estimators have their own sampling distributions.

14.2 Statistical Inference

What does this contain?

1. Point estimation
2. Hypothesis testing
3. Confidence Intervals

Example. Say we have 100,000 voters and we would like to know how many of them would vote for party A. Then let

$$X_i = \begin{cases} 1 & \text{if vote for A} \\ 0 & \text{otherwise} \end{cases}$$

then $S = \sum X_i$. Definitely, $S \sim \text{Bin}(n, p)$ where $P(X_i = 1) = p$. Then, $\mathbb{E}[S] = np \Rightarrow \mathbb{E}\left[\frac{S}{n}\right] = p$. Therefore let $\hat{p} = \frac{S}{n}$. So we will have an estimate of p given S and n . How about the error?

We know that $\mathbb{V}(S) = np(1-p)$. Now $X_1, \dots, X_n \stackrel{iid}{\sim} \text{Ber}(p)$. Using CLT,

$$\frac{\sqrt{n}(\bar{X} - \mu)}{\sigma} \stackrel{\text{approx}}{\sim} N(0, 1)$$

But $\bar{X} = \frac{S}{n} = \hat{p}$. Hence,

$$\frac{\sqrt{n}(\hat{p} - p)}{p(1-p)} \sim N(0, 1) \Leftrightarrow \hat{p} \stackrel{\text{approx}}{\approx} N\left(p, \frac{(p(1-p))^2}{n}\right)$$

Thus we take our error margin to be $\frac{\hat{p}(1-\hat{p})}{\sqrt{n}}$.